



Getting Set Up

A TERADATA COMPANY

Your environment (Tuesday, September 5)



- Each of you has your own Amazon Elastic MapReduce cluster with a unique IP address
- All are identical and have Hadoop, Spark, Hive, Pig, Python, and R installed on them

Name	Email		IP Address
Moore, David	david.moore2@elevate.com		34.201.251.109
Das, Debashis	debashis.das1@elevate.com		52.91.238.222
VonDohlen, Eric	eric.vondohlen@elevate.com		54.159.15.180
Shadrin, Ian	ian.shadrin@elevate.com		34.227.173.174
Seto, Jason	jason.seto@elevate.com		54.221.189.104
Baker, Jed	jed.baker@elevate.com		54.235.230.158
Chen, Lorraine	lorraine.chen@elevate.com		54.173.146.178
Domancic, Mirna	mirna.domancic@elevate.com		184.72.115.225
Hoelkher, Nick	nick.hoelkher@elevate.com		54.85.109.192
Gunasekar, Vikram	vikram.gunasekar@elevate.com		34.207.109.240
Zang, Xiaojiao	xiaojiao.zang@elevate.com		34.228.40.2
Chen, Yonghui	yonghui.che@elevate.com		34.207.65.85
Ahuja, Vijay	vahuja@elevate.com		34.201.126.127
Carl Howe	carl.howe@thinkbiganalytics.com		54.174.231.137

On a cluster



A TERADATA COMPANY

On completion, point your browser at
`http://your-ip-address-from-the-worksheet`

For example, Carl would type:

`http://52.86.53.106`

And he should see what's on the next page

Amazon Linux AMI Test Page

This page is used to test the proper operation of the Apache HTTP server after it has been installed. If you can read this page, it means that the web server installed at this site is working properly, but has not yet been configured.

If you are a member of the general public:

The fact that you are seeing this page indicates that the website you just visited is either experiencing problems, or is undergoing routine maintenance.

If you would like to let the administrators of this website know that you've seen this page instead of the page you expected, you should send them e-mail. In general, mail sent to the name 'webmaster' and directed to the website's domain should reach the appropriate person.

For example, if you experienced problems while visiting www.example.com, you should send e-mail to 'webmaster@example.com'.

The [Amazon Linux AMI](#) is a supported and maintained Linux image provided by [Amazon Web Services](#) for use on [Amazon Elastic Compute Cloud \(Amazon EC2\)](#). It is designed to provide a stable, secure, and high performance execution environment for applications running on [Amazon EC2](#). It also includes packages that enable easy integration with [AWS](#), including launch configuration tools and many popular AWS libraries and tools. [Amazon Web Services](#) provides ongoing security and maintenance updates to all instances running the [Amazon Linux AMI](#). The [Amazon Linux AMI](#) is provided at no additional charge to [Amazon EC2](#) users.

If you are the website administrator:

You may now add content to the directory `/var/www/html/`. Note that until you do so, people visiting your website will see this page, and not your content. To prevent this page from ever being used, follow the instructions in the file `/etc/httpd/conf.d/welcome.conf`.

You are free to use the images below on Apache and Amazon Linux AMI powered HTTP servers. Thanks for using Apache and the Amazon Linux AMI!



Setting Up Keys

Setting up cryptographic keys for access

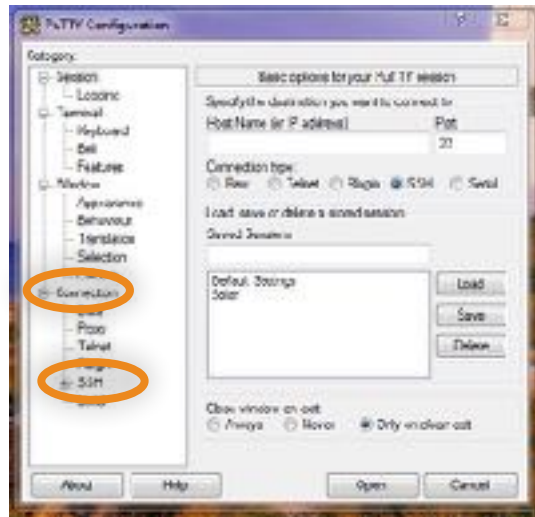


A TERADATA COMPANY

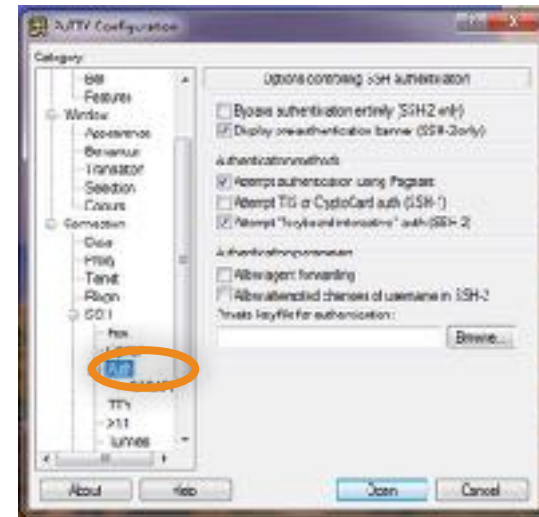
- Your Amazon EMR instance can only be used by someone with the appropriate cryptographic private key
- This course uses two different cryptographic private keys:
 - elv-key.pem for SSH clients
 - elv-key.ppk for PuTTY clients
- The following steps will walk you through installing a private cryptographic key into PuTTY

Setting up your PuTTY session



1. When you launch PuTTY you should see the following screen



2. Click “Connection > SSH > Auth” and you should see the screen below

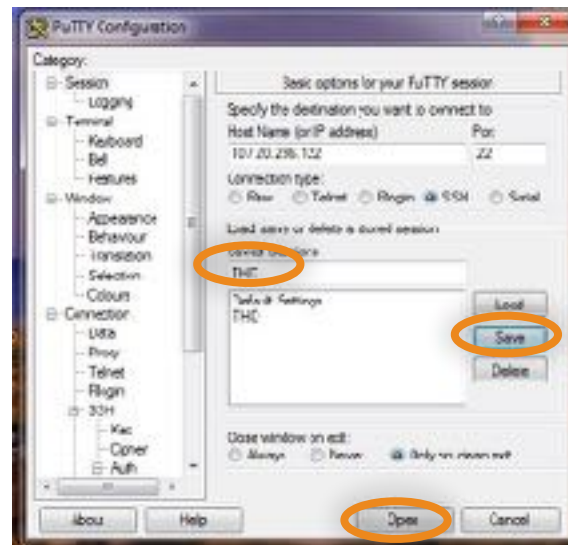



3. Browse to where you stored the private key distributed to you in email or on a thumb drive
4. Click Open for elv-key.ppk

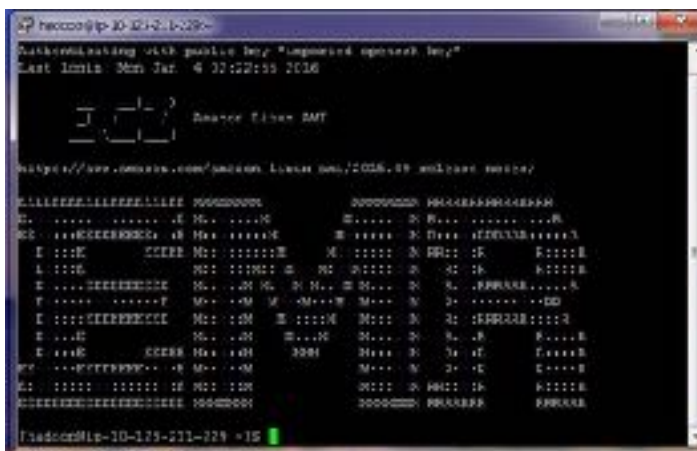


5. Return to the Session selection on the left and enter the IP address for your instance (this is Carl's)

6. Save your session as EMR and then click "Open"



7. Dismiss the security warning by clicking “Yes” and then type in your login name “hadoop”
8. You should now have a shell prompt on your



Congratulations!
You are up and running!





Instance Specs

General Purpose - Current Generation

t2.micro	1	Variable	1	EBS Only	\$0.013 per Hour
t2.small	1	Variable	2	EBS Only	\$0.026 per Hour
t2.medium	2	Variable	4	EBS Only	\$0.052 per Hour
t2.large	2	Variable	8	EBS Only	\$0.104 per Hour
m4.large	2	8.5	8	EBS Only	\$0.126 per Hour
m4.xlarge	4	13	16	EBS Only	\$0.252 per Hour
m4.2xlarge	8	26	32	EBS Only	\$0.504 per Hour
m4.4xlarge	16	53.5	64	EBS Only	\$1.008 per Hour
m4.10xlarge	40	124.5	160	EBS Only	\$2.52 per Hour
m3.medium	1	3	3.75	1 x 4 SSD	\$0.067 per Hour
m3.large	2	6.5	7.5	1 x 32 SSD	\$0.133 per Hour
m3.xlarge	4	13	15	2 x 40 SSD	\$0.266 per Hour
m3.2xlarge	8	26	30	2 x 80 SSD	\$0.532 per Hour